

Distributed Computing
Lab – I

K. Manoj Kumar
2022BCS0184.

PART – A:

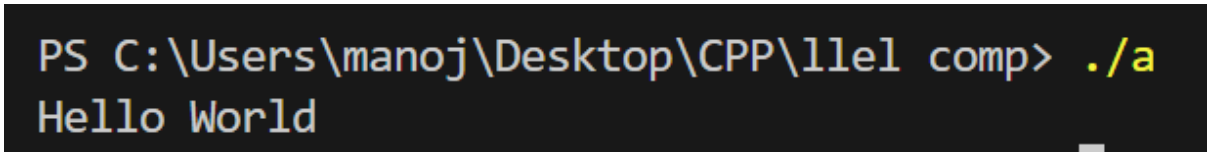
Q1. A). Write a simple C/C++ to print “Hello World”.

Code :

```
#include<iostream>

int main() {
    std::cout << "Hello World";
}
```

Output :



```
PS C:\Users\manoj\Desktop\CPP\11e1 comp> ./a
Hello World
```

B). Write an openMP program to print “Hello World” with parallel computation

Program :

```
#include<iostream>
#include<omp.h>

int main() {
    #pragma omp parallel
    {
        std::cout << "Hello World\n";
    }
}
```

```
}
```

Output :

```
PS C:\Users\manoj\Desktop\CPP\l1e1 comp> g++ -fopenmp q1.cpp
PS C:\Users\manoj\Desktop\CPP\l1e1 comp> ./a
Hello World
Hello World
Hello World
Hello World
Hello World
Hello World
Hello World
Hello World
Hello World
Hello World
Hello World
Hello World
Hello World
```

Q2. A). Write an openMP program to print Hello World with parallel computation along with the corresponding thread id.

Program :

```
#include<iostream>
#include<omp.h>
int main() {
    #pragma omp parallel
    {
        std::cout << "Hello World from Thread:" << omp_get_thread_num() <<
std::endl;
    }
}
```

Output :

```

PS C:\Users\manoj\Desktop\CPP\l1e1 comp> g++ -fopenmp q1.cpp
PS C:\Users\manoj\Desktop\CPP\l1e1 comp> ./a
Hello World from Thread:0
Hello World from Thread:7
Hello World from Thread:5
Hello World from Thread:2
Hello World from Thread:3
Hello World from Thread:9
Hello World from Thread:1
Hello World from Thread:1
Hello World from Thread:Hello World from Thread:11Hello World from Thread:10
Hello World from Thread:6
Hello World from Thread:8
4

```

B). Write an openMP program to print number of threads.

Program :

```

#include<iostream>
#include<omp.h>
int main() {
    #pragma omp parallel
    {
        #pragma omp single
        {
            std::cout << "No. of threads: " << omp_get_num_threads() << std::endl;
        }
    }
}

```

Output :

```

PS C:\Users\manoj\Desktop\CPP\l1e1 comp> ./a
No. of threads: 12

```

Q3. Write an openMP to print "Hello World" for k times using for loop and compute the execution time of sequential and parallel run.

Program:

```
#include<iostream>
```

```
#include<omp.h>
```

```
#include<chrono>
```

```
void sequential() {
```

```
    for(int i = 0; i < 10; i++) {
```

```
        std::cout << "Hello world\n";
```

```
    }
```

```
}
```

```
void parallel() {
```

```
    #pragma omp parallel for
```

```
        for(int i = 0; i < 10; i++)
```

```
        {
```

```
            std::cout << "Hello world from iteration : " << i << " thread: " <<
omp_get_thread_num() << std::endl;
```

```
        }
```

```
}
```

```
int main() {  
    auto s = std::chrono::high_resolution_clock::now();  
    sequential();  
    auto e = std::chrono::high_resolution_clock::now();  
    std::chrono::duration<double> dur = e - s;  
    std::cout << "Execution time for sequential:" << dur.count() << std::endl;  
    auto s_p = std::chrono::high_resolution_clock::now();  
    parallel();  
    auto e_p = std::chrono::high_resolution_clock::now();  
    dur = e_p - s_p;  
    std::cout << "Execution time for parallel:" << dur.count() << std::endl;  
  
}
```

Output:

```
Hello world
Hello world
Hello world
Hello world
Execution time for sequential:0.0023916
Hello world from iteration :Hello world from iteration :5thr
Hello world from iteration :8thread:8
Hello world from iteration :1thread:1
0thread:0
0thread:0
0thread:0
0thread:0
0thread:0
Hello world from iteration :7thread:Hello world from iterati
7
Hello world from iteration :2thread:2
Hello world from iteration :6thread:6
Hello world from iteration :3thread:3

Execution time for parallel:0.0043695
```

PART - B

Q4. Write an openMP program (C++) for matrix multiplication with serial and parallel computation and note the execution time of the same. Test the execution time with different dimension of the matrix (D). For example N=10, 100, 1000..... (Any Larger Number).

Program:

```
#include<bits/stdc++.h>
```

```
using namespace std;
```

```

void print(std::vector<std::vector<int>> arr) {
    int row = arr.size() ,column = arr[0].size();
    for(int i = 0;i < row;i++) {
        for(int j = 0;j < column;j++) {
            std::cout << arr[i][j] << " ";
        }
        std::cout << std::endl;
    }
}

```

```

void p_matmul(std::vector<std::vector<int>>
&arr1,std::vector<std::vector<int>> &arr2,std::vector<std::vector<int>>
&result,size_t dim) {
    #pragma omp parallel for collapse(3) shared(result)
    for(size_t i = 0;i < dim;i++) {
        for(size_t j = 0;j < dim;j++) {
            result[i][j] = 0;
            for(size_t k = 0;k < dim;k++) {

                result[i][j] += arr1[i][k] * arr2[k][j];
            }
        }
    }
}

```

```

void matmul(std::vector<std::vector<int>>
&arr1,std::vector<std::vector<int>> &arr2,std::vector<std::vector<int>>
&result,size_t dim) {

```

```

for(size_t i = 0; i < dim; i++) {
    for(size_t j = 0; j < dim; j++) {
        result[i][j] = 0;
        for(size_t k = 0; k < dim; k++) {

            result[i][j] += arr1[i][k] * arr2[k][j];

        }
    }
}

```

```

void fillMatrix(vector<vector<int>> &arr, size_t dimension) {
    std::random_device rd;
    unsigned seed =
std::chrono::system_clock::now().time_since_epoch().count();
    std::mt19937 gen(rd());
    std::uniform_int_distribution<> uni(0, 100);
    for(int i = 0; i < dimension; i++) {
        for(int j = 0; j < dimension; j++) {
            arr[i][j] = uni(gen);
        }
    }
}

```

```

void run(const size_t &dim) {

```



```

std::vector<std::vector<int>> arr1 (dim,std::vector<int>(dim));
std::vector<std::vector<int>> arr2 (dim,std::vector<int>(dim));
std::vector<std::vector<int>> result(dim,std::vector<int>(dim));
fillMatrix(arr1,dim);
fillMatrix(arr2,dim);
auto start = std::chrono::high_resolution_clock::now();
matmul(arr1,arr2,result,dim);
auto end = std::chrono::high_resolution_clock::now();
std::chrono::duration<double> duration = end - start;

auto start_p = std::chrono::high_resolution_clock::now();
p_matmul(arr1,arr2,result,dim);
auto end_p = std::chrono::high_resolution_clock::now();
std::chrono::duration<double> duration_p = end_p - start_p;

std::cout << "\nExecution time for Multiplication(serial) of matrix dimension
" << dim << " is : " << duration.count() ;

std::cout << "\nExecution time for Multiplication(parallel) : of matrix
dimension " << dim << " is : " << duration_p.count() ;

}

int main() {

for(int i : {10,100,1000}) {
    run(i);

```

```

    }

    return 0;
}

```

Output:

```

PS C:\Users\manoj\Desktop\CPP\l1e1 comp> g++ -fopenmp mat_mul.cpp
PS C:\Users\manoj\Desktop\CPP\l1e1 comp> ./a

Execution time for Multiplication(serial) of matrix dimension 10 is :7.1e-06
Execution time for Multiplication(parallel) : of matrix dimension 10is : 0.0023208
Execution time for Multiplication(serial) of matrix dimension 100 is :0.0065619
Execution time for Multiplication(parallel) : of matrix dimension 100is : 0.0027097
Execution time for Multiplication(serial) of matrix dimension 1000 is :5.22352
Execution time for Multiplication(parallel) : of matrix dimension 1000is : 1.90577
PS C:\Users\manoj\Desktop\CPP\l1e1 comp>

```

Q5. Write an openMP program (C++) for the following sorting algorithms with serial and parallel computation and note the execution time of the same. Analyse the output with different high N values, where N is the number of elements to be sorted. (Like above table) a) Merge sort b) Quick sort

Program(Merge Sort) :

```
#include<bits/stdc++.h>
```

```
using namespace std;
```

```
#define THRESHOLD 50
```

```
void merge(vector<int> &arr, int left, int mid, int right) {
```

```
    int l = mid - left + 1;
```

```
    int r = right - mid;
```

```
    vector<int> left_array(l);
```

```
    vector<int> right_array(r);
```

```

for (int i = 0; i < l; i++) {
    left_array[i] = arr[left + i];
}
for (int i = 0; i < r; i++) {
    right_array[i] = arr[mid + 1 + i];
}

```

```

int i = 0, j = 0, k = left;
while (i < l && j < r) {
    if (left_array[i] <= right_array[j]) {
        arr[k] = left_array[i];
        i++;
    } else {
        arr[k] = right_array[j];
        j++;
    }
    k++;
}

```

```

while (i < l) {
    arr[k] = left_array[i];
    i++;
    k++;
}

```

```

while (j < r) {

```

```

        arr[k] = right_array[j];
        j++;
        k++;
    }
}

void print(vector<int>& arr) {
    std::cout << "[";
    for (int i : arr) {
        cout << i << ", ";
    }
    std::cout << "]";
}

```

```

void mergeSort(vector<int> &arr, const int &left, const int &right) {
    if (left >= right)
        return;
    int mid = left + (right - left) / 2;
    mergeSort(arr, left, mid);
    mergeSort(arr, mid + 1, right);
    merge(arr, left, mid, right);
}

```

```

void mergeParallelSort(vector<int> &arr, const int &left, const int &right) {
    if(left >= right) {
        return;
    }
}

```

```

int mid = left + (right - left) / 2;
if(right - left < THRESHOLD) {
    mergeSort(arr, left, right);
    return;
}
#pragma omp parallel
{
    #pragma omp single
    {
        #pragma omp taskgroup
        {
            #pragma omp task shared(arr)
            mergeParallelSort(arr, left, mid);

            #pragma omp task shared(arr)
            mergeParallelSort(arr, mid + 1, right);
        }
    }

}

#pragma omp taskwait
merge(arr, left, mid, right);
}

void fill_vector(vector<int> &arr) {
    std::random_device rd;

```

```

    unsigned seed =
std::chrono::system_clock::now().time_since_epoch().count();

    std::mt19937 gen(rd());
    std::uniform_int_distribution<> uni(0, 100);
    for(int i = 0; i < arr.size(); i++) {
        arr[i] = uni(gen);
    }
}

```

```

void run(const size_t &array_length) {
    vector<int> arr(array_length);

    fill_vector(arr);
    auto s = std::chrono::high_resolution_clock::now();
    mergeSort(arr, 0, array_length - 1);
    auto e = std::chrono::high_resolution_clock::now();
    std::chrono::duration<double> run_time = e - s;

    std::cout << "size [" << std::setw(4) << std::left << array_length << "] runtime
for serial merge sort " << run_time.count() << "s" << std::endl;

```

```

    fill_vector(arr);
    auto ps = std::chrono::high_resolution_clock::now();
    mergeParallelSort(arr, 0, array_length - 1);
    auto pe = std::chrono::high_resolution_clock::now();
    run_time = pe - ps;

    std::cout << "size [" << std::setw(4) << array_length << "] runtime for parall
merge sort " << run_time.count() << "s" << std::endl;

```

```
}
```

```
int main() {
```

```
    for(const auto& i : {10,100,1000}) {
```

```
        run(i);
```

```
    }
```

```
}
```

Output:

```
PS C:\Users\manoj\Desktop\CPP\l1e1 comp> g++ -fopenmp .\merge_sort.cpp
PS C:\Users\manoj\Desktop\CPP\l1e1 comp> ./a
size [10 ] runtime for serial merge sort 6.3e-06s
size [10 ] runtime for parall merge sort 3.8e-06s
size [100 ] runtime for serial merge sort 4.69e-05s
size [100 ] runtime for parall merge sort 0.0036892s
size [1000] runtime for serial merge sort 0.000344s
size [1000] runtime for parall merge sort 0.0007356s
```

Program (Quick Sort) :

```
#include<bits/stdc++.h>
```

```
using namespace std;
```

```
int partition (vector<int>& arr,int left,int right) {
```

```
    int piv = arr[right];
```

```
    int l = left - 1;
```

```
    for(int i = left;i < right;i++) {
```

```
        if(arr[i] < piv) {
```

```
            l++;
```

```

        swap(arr[l],arr[i]);
    }
}
swap(arr[l + 1],arr[right]);
return l + 1;
}

```

```

void quick(vector<int>& arr,int left,int right) {
    if(left < right) {
        int pivot = partition(arr,left,right);
        quick(arr,left,pivot - 1);
        quick(arr,pivot + 1,right);
    }
}

```

```

void print(vector<int>& arr) {
    std::cout << "[";
    for (int i : arr) {
        cout << i << ", ";
    }
    std::cout << "]";
}

```

```

void p_quick(vector<int>& arr,int left,int right) {
    if (left < right) {
        int pivot = partition(arr, left, right);
    }
}

```



```

// Parallelize recursive calls if the partition is large enough
#pragma omp parallel sections if(right - left > 50) // Limit small partitions
from spawning too many threads
{
    #pragma omp section
    {
        p_quick(arr, left, pivot - 1);
    }
    #pragma omp section
    {
        p_quick(arr, pivot + 1, right);
    }
}
}
}

```

```

void randomize(vector<int> &arr) {
    std::random_device rd;
    int n = arr.size();
    unsigned seed =
std::chrono::system_clock::now().time_since_epoch().count();
    std::mt19937 gen(rd());
    std::uniform_int_distribution<> uni(INT16_MIN, INT16_MAX);
    for(int i = 0; i < n; i++) {
        arr[i] = uni(gen);
    }
}

```

```
}
```

```
void run(const size_t &length) {
```

```
    vector<int> arr(length);
```

```
    randomize(arr);
```

```
    auto start = std::chrono::high_resolution_clock::now();
```

```
    quick(arr,0,arr.size() - 1);
```

```
    auto end = std::chrono::high_resolution_clock::now();
```

```
    std::chrono::duration<double> duration = end - start;
```

```
    std::cout << "\nexecution time for serial of array length "<< length << " is : "  
<< duration.count() << endl;
```

```
    randomize(arr);
```

```
    auto start_p = std::chrono::high_resolution_clock::now();
```

```
    #pragma omp parallel
```

```
{
```

```
        #pragma omp single
```

```
{
```

```
            p_quick(arr, 0, arr.size() - 1);
```

```
}
```

```
}
```

```
    auto end_p = std::chrono::high_resolution_clock::now();
```

```
    std::chrono::duration<double> duration_p = end_p - start_p;
```

```

        std::cout << "\nExecution time for parallel of array length "<< length << " is :"  

        << duration_p.count() << endl;  

    }

```

```

int main() {
    for(int len : {10,100,1000}) {
        run(len);
    }
    return 0;
}

```

Output:

```

execution time for serial of array length 10 is : 7e-07
Execution time for parallel of array length 10 is :0.0126136
execution time for serial of array length 100 is : 6.7e-06
Execution time for parallel of array length 100 is :0.000362
execution time for serial of array length 1000 is : 0.0001029
Execution time for parallel of array length 1000 is :0.0021814

```